

```

library(car)

## Loading required package: carData

library(reshape2)
library(ggplot2)
library(MASS)
library(interactions)

source("clean_data.R")

df <- remove_cols(df, c("Color", "Model"))

# Remove columns with only one observation and affected rows
res <- convert_categorical(df, categorical)
design <- as.data.frame(res$dummy)

singles <- c()
bad_idx <- c()
for (col in colnames(design)) {
  if (sum(design[, col] != 0) <= 1) {
    singles <- c(singles, col)
    bad_idx <- c(bad_idx, which(design[, col] != 0))
  }
}
singles

## [1] "MakeFiat" "MakeLexus" "Fuel.TypePetrol + CNG"
## [4] "LocationDak. Kannada" "LocationFaizabad" "LocationGorakhpur"
## [7] "LocationPurnea" "LocationRohtak" "LocationRudrapur"
## [10] "LocationSamastipur" "LocationValsad"


bad_idx

## [1] 662 1009 1169 1077 779 569 728 245 510 264 162
df <- df[-bad_idx, ]

```

## Box Cox

```
bc <- boxcox(Price ~., data = df)
```



```
bc$x[which.max(bc$y)]
```

```
## [1] -0.06060606
```

For the sake of interpretability, we will use a log transformation, since  $\lambda \approx 0$ .

## Iteratively Remove Multicollinear Regressors

```
x <- df
model <- lm(log(Price) ~., data = x)
removed <- c()

finished <- F
while(!finished) {
  temp <- car::vif(model)[, "GVIF^(1/(2*Df))"]
  worst <- names(which.max(temp))
  if (length(temp) > 0 && temp[worst] > sqrt(5)) {
    x <- remove_cols(x, c(worst))
    model <- lm(log(Price) ~., data = x)

    removed <- c(removed, worst)
  } else {
    finished <- T
  }
}

removed

## [1] "Max.Torque.Value" "Max.Power.RPM" "Max.Power.Value"
## [4] "Engine" "Fuel.Tank.Capacity" "Height"
which(abs(rstudent(model)) > 4)

## 44 366 492 1536 1681
```

```
## 39 268 345 1071 1171
```

After inspecting these data points, we do not find a good reason to remove them (i.e., they are not clerical errors).

```
vif(model, type = "predictor")
```



```
## GVIFs computed for predictors
```


```
##           GVIF Df GVIF^(1/(2*Df)) Interacts With
## Make          76.706634 25         1.090678      --
## Year           2.038829  1         1.427876      --
## Kilometer      2.089964  1         1.445671      --
## Fuel.Type      5.026586  2         1.497333      --
## Transmission   1.843543  1         1.357771      --
## Location      13.438914 63         1.020834      --
## Owner          1.628105  3         1.084627      --
## Seller.Type    1.329867  2         1.073871      --
## Max.Torque.RPM 4.220839  1         2.054468      --
## Drivetrain     3.872635  2         1.402819      --
## Length         4.511737  1         2.124085      --
## Width          4.596043  1         2.143838      --
## Seating.Capacity 2.506336 1         1.583141      --
```

```
## Make          Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Year          Make, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Kilometer     Make, Year, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Fuel.Type     Make, Year, Kilometer, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Transmission  Make, Year, Kilometer, Fuel.Type, Location, Owner, Seller.Type, Max.Torque.L
## Location      Make, Year, Kilometer, Fuel.Type, Transmission, Owner, Seller.Type, Max.Torque.L
## Owner         Make, Year, Kilometer, Fuel.Type, Transmission, Location, Seller.Type, Max.Torque.L
## Seller.Type   Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Max.Torque.L
## Max.Torque.RPM Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Drivetrain    Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Length        Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Width         Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
## Seating.Capacity Make, Year, Kilometer, Fuel.Type, Transmission, Location, Owner, Seller.Type, Max.Torque.L
```

We get that  $GVIF^{1/(2 \cdot Df)} < 2.236068 \approx \sqrt{5}$  for all regressors, and no interactions are indicated, so there is likely no multicollinearity.

```
plot(model)
```





## Backwards Stepwise Search

We do a backwards search since our model already conforms to the linear assumptions and is performing well. We simply wish to reduce the model size now.

```
reduced_model <- step(model, direction = "backward", data = x, trace = 0, k = log(nrow(x)))
old <- names(model$coefficients)
new <- names(reduced_model$coefficients)

old[!(old %in% new)]
```

```
## [1] "LocationAhmedabad"      "LocationAllahabad"      "LocationAmbala Cantt"
## [4] "LocationAmritsar"       "LocationAurangabad"     "LocationBangalore"
## [7] "LocationBhopal"        "LocationBhubaneswar"    "LocationBulandshahar"
## [10] "LocationChandigarh"     "LocationChennai"        "LocationCoimbatore"
## [13] "LocationDehradun"      "LocationDelhi"          "LocationDharwad"
## [16] "LocationErnakulam"     "LocationFaridabad"      "LocationGhaziabad"
## [19] "LocationGoa"           "LocationGurgaon"        "LocationGuwahati"
## [22] "LocationHaldwani"       "LocationHyderabad"      "LocationIndore"
## [25] "LocationJaipur"        "LocationJalandhar"      "LocationJamshedpur"
## [28] "LocationKanpur"        "LocationKarnal"         "LocationKharar"
## [31] "LocationKheda"         "LocationKolkata"        "LocationKollam"
## [34] "LocationKota"          "LocationLucknow"        "LocationLudhiana"
## [37] "LocationMangalore"     "LocationMeerut"         "LocationMirzapur"
## [40] "LocationMohali"        "LocationMumbai"         "LocationMuzaffarpur"
## [43] "LocationMysore"        "LocationNagpur"         "LocationNashik"
## [46] "LocationNavi Mumbai"   "LocationNoida"          "LocationPanchkula"
## [49] "LocationPatna"         "LocationPune"           "LocationRaipur"
## [52] "LocationRanchi"        "LocationRanga Reddy"    "LocationRoorkee"
## [55] "LocationSalem"         "LocationSurat"          "LocationThane"
## [58] "LocationUdupi"         "LocationVadodara"       "LocationVaranasi"
## [61] "LocationWarangal"      "LocationYamunanagar"    "LocationZirakpur"
## [64] "OwnerSecond"           "OwnerThird"             "OwnerUnRegistered Car"
## [67] "Seller.TypeCorporate"  "Seller.TypeIndividual"  "Seating.Capacity"
```

Location, seller type, owner type, and seating capacity were removed by BIC.

## Inspect Model Coefficients

```
summary(reduced_model)
```

```
##
## Call:
## lm(formula = log(Price) ~ Make + Year + Kilometer + Fuel.Type +
##     Transmission + Max.Torque.RPM + Drivetrain + Length + Width,
##     data = x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04626 -0.11894  0.00895  0.12150  1.12776
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -2.108e+02  4.959e+00 -42.513  < 2e-16 ***
## MakeBMW       -8.289e-03  4.487e-02  -0.185  0.85345
## MakeChevrolet -7.614e-01  1.036e-01  -7.353  3.28e-13 ***
## MakeDatsun    -9.462e-01  7.807e-02 -12.120  < 2e-16 ***
## MakeFord     -5.729e-01  4.202e-02 -13.635  < 2e-16 ***
## MakeHonda    -5.122e-01  3.344e-02 -15.318  < 2e-16 ***
```

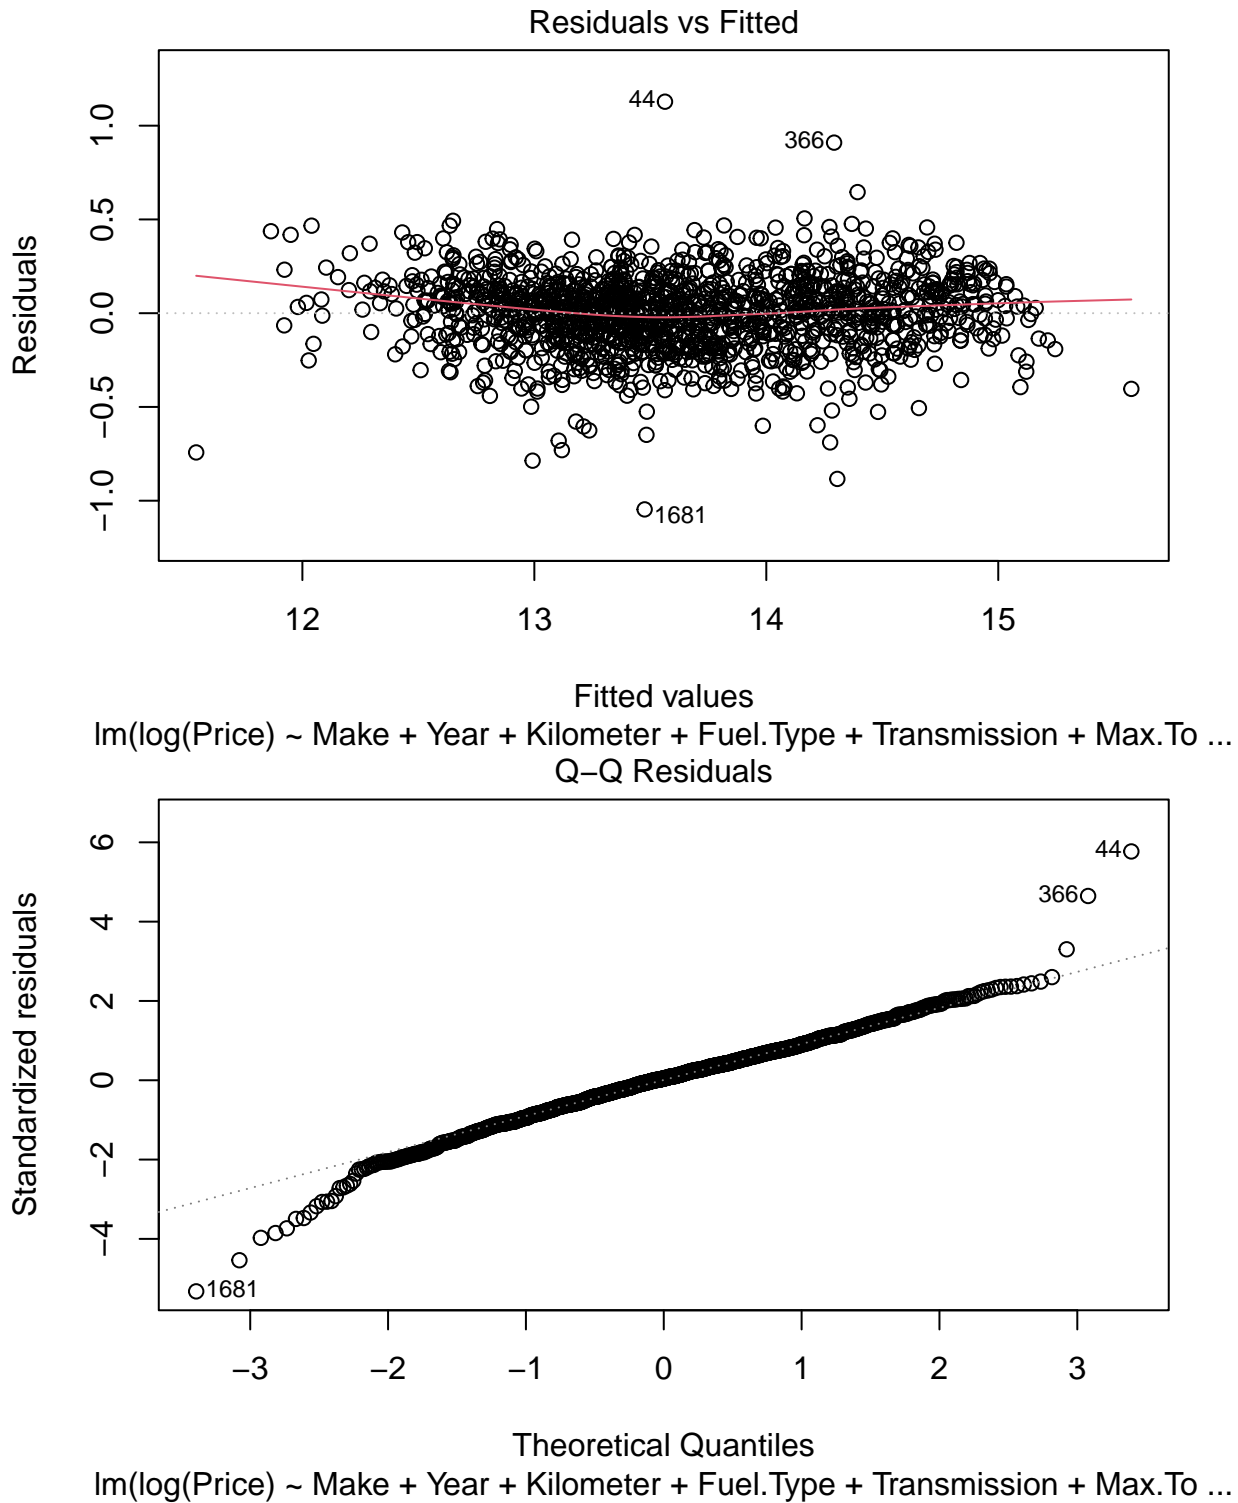
```
## MakeHyundai      -4.787e-01  2.967e-02 -16.136 < 2e-16 ***
## MakeIsuzu        -3.579e-01  1.445e-01  -2.477  0.01336 *
## MakeJaguar        2.577e-02  8.764e-02   0.294  0.76879
## MakeJeep         -3.124e-01  5.672e-02  -5.508  4.31e-08 ***
## MakeKia          -4.186e-01  4.935e-02  -8.482 < 2e-16 ***
## MakeLand Rover    1.135e-01  1.180e-01   0.962  0.33613
## MakeMahindra     -6.140e-01  3.531e-02 -17.389 < 2e-16 ***
## MakeMaruti Suzuki -5.372e-01  3.052e-02 -17.602 < 2e-16 ***
## MakeMercedes-Benz 9.442e-02  3.692e-02   2.558  0.01065 *
## MakeMG           -5.046e-01  6.084e-02  -8.295  2.52e-16 ***
## MakeMINI          6.884e-01  8.562e-02   8.041  1.88e-15 ***
## MakeMitsubishi   -2.069e-01  1.185e-01  -1.746  0.08110 .
## MakeNissan       -7.701e-01  5.774e-02 -13.338 < 2e-16 ***
## MakeRenault      -7.636e-01  4.273e-02 -17.869 < 2e-16 ***
## MakeSkoda        -4.656e-01  4.166e-02 -11.176 < 2e-16 ***
## MakeSsangyong    -8.916e-01  1.188e-01  -7.502  1.11e-13 ***
## MakeTata         -7.320e-01  4.051e-02 -18.068 < 2e-16 ***
## MakeToyota       -2.752e-01  3.299e-02  -8.341 < 2e-16 ***
## MakeVolkswagen   -5.016e-01  4.107e-02 -12.213 < 2e-16 ***
## MakeVolvo        3.521e-02  6.401e-02   0.550  0.58236
## Year             1.094e-01  2.470e-03  44.282 < 2e-16 ***
## Kilometer       -1.624e-06  2.597e-07  -6.253  5.35e-10 ***
## Fuel.TypeDiesel  -1.320e-01  4.157e-02  -3.175  0.00153 **
## Fuel.TypePetrol  -1.518e-01  3.794e-02  -4.000  6.67e-05 ***
## TransmissionManual -1.660e-01  1.395e-02 -11.906 < 2e-16 ***
## Max.Torque.RPM   -4.436e-05  8.320e-06  -5.331  1.14e-07 ***
## DrivetrainFWD    -2.586e-01  2.475e-02 -10.449 < 2e-16 ***
## DrivetrainRWD    -1.703e-01  2.968e-02  -5.738  1.17e-08 ***
## Length           5.894e-04  2.775e-05  21.239 < 2e-16 ***
## Width            1.444e-03  1.208e-04  11.951 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1986 on 1405 degrees of freedom
## Multiple R-squared:  0.9207, Adjusted R-squared:  0.9187
## F-statistic: 465.9 on 35 and 1405 DF, p-value: < 2.2e-16
```

```
anova(reduced_model)
```



```
## Analysis of Variance Table
##
## Response: log(Price)
##              Df Sum Sq Mean Sq  F value    Pr(>F)
## Make          25 354.44   14.178   359.289 < 2.2e-16 ***
## Year           1 178.52  178.525  4524.169 < 2.2e-16 ***
## Kilometer      1   0.98   0.982    24.874 6.884e-07 ***
## Fuel.Type      2  19.23   9.615   243.668 < 2.2e-16 ***
## Transmission   1  20.31  20.313   514.760 < 2.2e-16 ***
## Max.Torque.RPM 1   1.14   1.139    28.868 9.060e-08 ***
## Drivetrain     2   6.53   3.267    82.792 < 2.2e-16 ***
## Length         1  56.67  56.670  1436.142 < 2.2e-16 ***
## Width          1   5.64   5.636   142.825 < 2.2e-16 ***
## Residuals     1405  55.44   0.039
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Final Model Verification

```
plot(reduced_model)
```







### Normalized Coefficients

```
# removed by BIC. remember multicollinear rows are already removed
normalized <- remove_cols(x, c("Location", "Seating.Capacity", "Seller.Type", "Owner"))
```

```

for (col in colnames(normalized)) {
  if (col != "Price" && !(col %in% categorical)) {
    normalized[, col] <- (normalized[, col] - mean(normalized[, col])) / sd(normalized[, col])
  }
}

normalized_model <- lm(log(Price) ~ ., data = normalized)
summary(normalized_model)

```

```

##
## Call:
## lm(formula = log(Price) ~ ., data = normalized)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04626 -0.11894  0.00895  0.12150  1.12776
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    14.543827   0.045612  318.858 < 2e-16 ***
## MakeBMW        -0.008289   0.044869   -0.185  0.85345
## MakeChevrolet  -0.761421   0.103550   -7.353 3.28e-13 ***
## MakeDatsun     -0.946201   0.078066  -12.120 < 2e-16 ***
## MakeFord       -0.572886   0.042016  -13.635 < 2e-16 ***
## MakeHonda      -0.512179   0.033437  -15.318 < 2e-16 ***
## MakeHyundai    -0.478712   0.029667  -16.136 < 2e-16 ***
## MakeIsuzu      -0.357869   0.144462   -2.477  0.01336 *
## MakeJaguar      0.025768   0.087639    0.294  0.76879
## MakeJeep       -0.312446   0.056724   -5.508 4.31e-08 ***
## MakeKia        -0.418627   0.049353   -8.482 < 2e-16 ***
## MakeLand Rover  0.113527   0.117990    0.962  0.33613
## MakeMahindra   -0.613960   0.035308  -17.389 < 2e-16 ***
## MakeMaruti Suzuki -0.537170   0.030517  -17.602 < 2e-16 ***
## MakeMercedes-Benz 0.094425   0.036921    2.558  0.01065 *
## MakeMG         -0.504646   0.060838   -8.295 2.52e-16 ***
## MakeMINI        0.688415   0.085617    8.041 1.88e-15 ***
## MakeMitsubishi -0.206925   0.118540   -1.746  0.08110 .
## MakeNissan      -0.770128   0.057739  -13.338 < 2e-16 ***
## MakeRenault     -0.763589   0.042732  -17.869 < 2e-16 ***
## MakeSkoda       -0.465623   0.041663  -11.176 < 2e-16 ***
## MakeSsangyong   -0.891608   0.118844   -7.502 1.11e-13 ***
## MakeTata        -0.732018   0.040515  -18.068 < 2e-16 ***
## MakeToyota      -0.275157   0.032989   -8.341 < 2e-16 ***
## MakeVolkswagen  -0.501595   0.041071  -12.213 < 2e-16 ***
## MakeVolvo       0.035208   0.064006    0.550  0.58236
## Year            0.311815   0.007042   44.282 < 2e-16 ***
## Kilometer      -0.044526   0.007121   -6.253 5.35e-10 ***
## Fuel.TypeDiesel -0.132002   0.041572   -3.175  0.00153 **
## Fuel.TypePetrol -0.151767   0.037944   -4.000 6.67e-05 ***
## TransmissionManual -0.166038   0.013946  -11.906 < 2e-16 ***
## Max.Torque.RPM  -0.055706   0.010449   -5.331 1.14e-07 ***
## DrivetrainFWD   -0.258640   0.024754  -10.449 < 2e-16 ***
## DrivetrainRWD   -0.170290   0.029677   -5.738 1.17e-08 ***

```

```
## Length          0.218156    0.010272   21.239   < 2e-16 ***
## Width           0.130065    0.010883   11.951   < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1986 on 1405 degrees of freedom
## Multiple R-squared:  0.9207, Adjusted R-squared:  0.9187
## F-statistic: 465.9 on 35 and 1405 DF,  p-value: < 2.2e-16
sort(abs(normalized_model$coefficients), decreasing = T)
```

```
##      (Intercept)      MakeDatsun      MakeSsangyong      MakeNissan
##      14.543826960      0.946201111      0.891608072      0.770128355
##      MakeRenault      MakeChevrolet      MakeTata      MakeMINI
##      0.763589078      0.761420626      0.732018001      0.688415224
##      MakeMahindra      MakeFord      MakeMaruti Suzuki      MakeHonda
##      0.613960112      0.572885621      0.537170191      0.512179318
##      MakeMG      MakeVolkswagen      MakeHyundai      MakeSkoda
##      0.504645860      0.501595422      0.478711619      0.465623344
##      MakeKia      MakeIsuzu      MakeJeep      Year
##      0.418626911      0.357868994      0.312445876      0.311815408
##      MakeToyota      DrivetrainFWD      Length      MakeMitsubishi
##      0.275156539      0.258640046      0.218155692      0.206925140
##      DrivetrainRWD TransmissionManual      Fuel.TypePetrol      Fuel.TypeDiesel
##      0.170289761      0.166038313      0.151767051      0.132001751
##      Width      MakeLand Rover      MakeMercedes-Benz      Max.Torque.RPM
##      0.130064979      0.113527147      0.094424967      0.055705512
##      Kilometer      MakeVolvo      MakeJaguar      MakeBMW
##      0.044525748      0.035207673      0.025767531      0.008289357
```

## Confidence Intervals

```
confint(normalized_model, level = 0.95)

##              2.5 %      97.5 %
## (Intercept)  14.45435143 14.63330249
## MakeBMW      -0.09630597  0.07972726
## MakeChevrolet -0.96455003 -0.55829122
## MakeDatsun   -1.09934049 -0.79306173
## MakeFord     -0.65530709 -0.49046415
## MakeHonda    -0.57777174 -0.44658689
## MakeHyundai  -0.53690767 -0.42051557
## MakeIsuzu    -0.64125440 -0.07448359
## MakeJaguar   -0.14615075  0.19768581
## MakeJeep     -0.42371850 -0.20117325
## MakeKia      -0.51543960 -0.32181422
## MakeLand Rover -0.11792897  0.34498326
## MakeMahindra -0.68322180 -0.54469842
## MakeMaruti Suzuki -0.59703421 -0.47730618
## MakeMercedes-Benz 0.02199928  0.16685065
## MakeMG       -0.62398874 -0.38530298
## MakeMINI      0.52046454  0.85636591
## MakeMitsubishi -0.43945862  0.02560835
## MakeNissan    -0.88339267 -0.65686404
## MakeRenault   -0.84741387 -0.67976428
```

## MakeSkoda	-0.54735131	-0.38389537
## MakeSsangyong	-1.12473817	-0.65847798
## MakeTata	-0.81149433	-0.65254168
## MakeToyota	-0.33986892	-0.21044416
## MakeVolkswagen	-0.58216290	-0.42102794
## MakeVolvo	-0.09034961	0.16076496
## Year	0.29800228	0.32562854
## Kilometer	-0.05849508	-0.03055642
## Fuel.TypeDiesel	-0.21355073	-0.05045278
## Fuel.TypePetrol	-0.22620082	-0.07733328
## TransmissionManual	-0.19339565	-0.13868098
## Max.Torque.RPM	-0.07620280	-0.03520823
## DrivetrainFWD	-0.30719809	-0.21008200
## DrivetrainRWD	-0.22850512	-0.11207440
## Length	0.19800639	0.23830499
## Width	0.10871583	0.15141413